



WRITE A POST



João Ventura

FOLLOW

Python enthusiast with 10 years of experience!

Building a basic HTTP Server from scratch in Python

Published Nov 22, 2020



In its essence, the modern web is just text going back and forth between clients and servers. As developers, we often use web frameworks to help us build strings to send to the clients. Web frameworks abstract us from the underlying "textual reality" by parsing the incoming http requests (which is just a string), call the corresponding function, and build a string response (mostly by using templates). Finally, clients parse those strings and do whatever they want from it.

This blog post shows how to build a barebones HTTP server from scratch and it is based on an exercise I gave to my MSc students. The only pre-requisite is a basic understanding of Python 3. If you want to implement this as we go along, you can grab the starting application [from this link](#). The final source code can be found in [this gist](#).

HTTP is just text

HTTP is the [protocol](#) that browsers use to retrieve and push information to servers. In its essence HTTP is just text that follows a certain pattern: on the first line you specify which resource you want, then it follows the headers, and then you have a blank line that separates the headers from the body of the message (if any). Here's how you could retrieve the about page of a website:

GET /about

By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



Enjoy this post?



And here's how you could send some form data to a web server, using the POST method:

```
POST /form.php HTTP/1.0
Content-Type: application/x-www-form-urlencoded
Content-Length: 21

name=John&surname=Doe
```

To prove that this is only text, you could simply copy-paste the text above and use something that allows you to send text over a network. Let's use telnet to get the google's about page:

```
$> telnet google.com 80
Trying 84.91.171.170...
Connected to google.com.
Escape character is '^]'.

(1)
GET /about/ HTTP/1.0

(2)
HTTP/1.0 200 OK
Vary: Accept-Encoding
Content-Type: text/html
Date: Thu, 09 Feb 2017 16:41:37 GMT
Expires: Thu, 09 Feb 2017 16:41:37 GMT
Cache-Control: private, max-age=0
Last-Modified: Thu, 08 Dec 2016 01:00:57 GMT
X-Content-Type-Options: nosniff
Server: sffe
X-XSS-Protection: 1; mode=block
Accept-Ranges: none
```

```
<!DOCTYPE html>
<html class="google mmfb" lang="en">
```

By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



Enjoy this post?



Connection closed by foreign host.

and send a POST message to <http://httpbin.org/post>:

```
$> telnet httpbin.org 80

Trying 54.175.219.8...
Connected to httpbin.org.
Escape character is '^]'.

(1)
POST /post HTTP/1.0
Content-Type: application/x-www-form-urlencoded
Content-Length: 21

name=John&surname=Doe

(2)
HTTP/1.1 200 OK
Server: nginx
Date: Thu, 09 Feb 2017 16:38:26 GMT
Content-Type: application/json
Content-Length: 328
Connection: close
Access-Control-Allow-Origin: *
Access-Control-Allow-Credentials: true
```

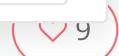
{
 "form": {
 "name": "John",
 "surname": "Doe"
 }
}

By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



Enjoy this post?



```
"Content-Length": "21",
"Content-Type": "application/x-www-form-urlencoded",
"Host": "httpbin.org"
},
"url": "http://httpbin.org/post"
}
Connection closed by foreign host.
```

You can see the HTTP requests (1) followed by the HTTP server responses in (2). Same pattern on requests and responses, and text everywhere! More information about HTTP on the excellent [High Performance Browser Networking](#) book.

Sending HTTP responses using sockets

If you are planning to implement network applications from scratch, you'll probably need to work with [network sockets](#). A socket is an abstraction provided by your operating system that allows you to send and receive bytes through a network. Here's a basic implementation of an HTTP server (you can get it [from this link](#)):

```
"""
Implements a simple HTTP/1.0 Server

"""

import socket

# Define socket host and port
SERVER_HOST = '0.0.0.0'
SERVER_PORT = 8000

# Create socket
server_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
server_socket.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
server_socket.bind((SERVER_HOST, SERVER_PORT))
server_socket.listen(1)
print('Listening on port %s ...' % SERVER_PORT)

while True:
    # Wa:
```

By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



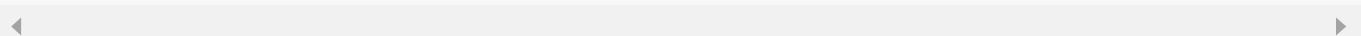
Enjoy this post?



```
request = client_connection.recv(1024).decode()
print(request)

# Send HTTP response
response = 'HTTP/1.0 200 OK\n\nHello World'
client_connection.sendall(response.encode())
client_connection.close()

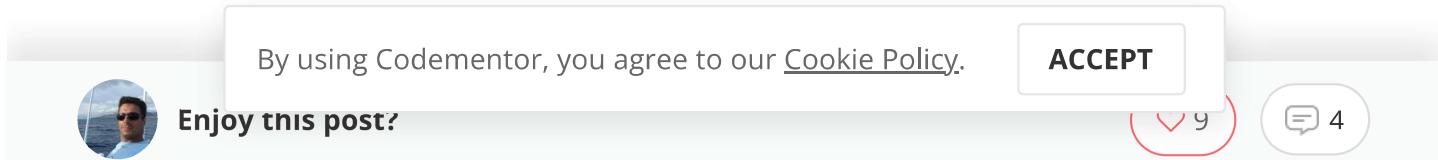
# Close socket
server_socket.close()
```



We start by defining the socket host and port. Then, we create the `server_socket` variable and set it to `AF_INET` (IPv4 address family) and `SOCK_STREAM` (TCP, basically).

The rest of the code is there to set up the socket to listen for requests on the given (host, port). Check the [Python docs on sockets](#) for more info.

The rest of the code is self-explanatory: wait for client connections, read the request string, send an HTTP-formatted string with *Hello World* on the response body and close the client connection. We do this forever (or until someone presses Ctrl+C). Open your browser on <http://localhost:8000/> and you should see the server's response:



As an exercise, change the *Hello World* to `<h1>Hello World</h1>` and see what happens. And did you see the `print(request)` in the server's source code? Here's what it outputs in the console:

```
Listening on port 8000 ...
```

```
GET / HTTP/1.1
Host: localhost:8080
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10.12; rv:51.0) Gecko/20100101 Firefox/51.0
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
Accept-Language: pt-PT,pt;q=0.8,en;q=0.5,en-US;q=0.3
Accept-Encoding: gzip, deflate
Connection: keep-alive
Upgrade-Insecure-Requests: 1
Cache-Control: max-age=0
```

Yes, it's the browser requesting the root page ("`/`") of the server..

Index.html

By default, when a browser requests the root of a server (using an HTTP request such as **GET / HTTP/1.0**), we should return the *index.html* page. Let's change the code inside the while to always return the contents of the *htdocs/index.html* file.

```
while True:
    # Wait for client connections
    (...)
```

By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



Enjoy this post?



```
# Get the content of htdocs/index.html
fin = open('htdocs/index.html')
content = fin.read()
fin.close()

# Send HTTP response
response = 'HTTP/1.0 200 OK\n\n' + content
client_connection.sendall(response.encode())
(...)
```

Basically, we read the contents of the file and add it to the *response* string as message body, instead of the previous *Hello World*. The *index.html* file is just a text file (inside the *htdocs* directory) with html content:

```
<html>
<head>
    <title>Hello World</title>
</head>
<body>
    <h1>Hello World!</h1>
    <p>Welcome to the index.html web page..</p>
    <p>Here's a link to <a href="ipsum.html">Ipsum</a></p>
</body>
</html>
```

Here's how it should look like in the browser:

By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



Enjoy this post?



You can click on the link as many times as you want, but your server will **always** return the contents of *index.html*. It was programmed to behave that way!

Return other pages

So far our server returns the *index.html* page but we should allow it to return other pages as well. Technically, it means that we must parse the first line of the HTTP request (which is something like *GET /ipsum.html HTTP/1.0*), open the intended file and return its contents. Here's the changes:

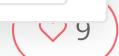
```
while True:  
    # Wait for client connections  
    (...)  
    # Get the client request  
    (...)  
  
    # Parse HTTP headers  
    headers = request.split('\n')  
    filename = headers[0].split()[1]  
  
    # Get the content of the file  
    if filename == '/':  
        filename = '/index.html'  
  
    fin = open('htdocs' + filename)
```

content By using Codementor, you agree to our [Cookie Policy](#).

ACCEPT



Enjoy this post?



```
# Send HTTP response
response = 'HTTP/1.0 200 OK\n\n' + content
client_connection.sendall(response.encode())
(...)
```

Basically, we extract the filename from the request string, open the file (we assume that all html files are inside the `htdocs` folder) and return its content. You can also check that we correctly return the `index.html` file when the clients ask for the root resource ('/').

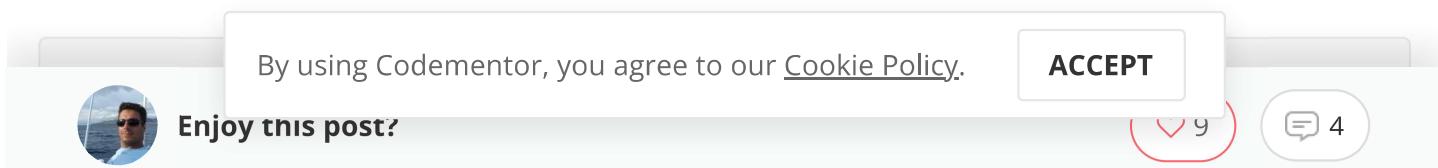
Here's the content of `htdocs/ipsum.html`:

```
<html>
<head>
    <title>Ipsum</title>
</head>
<body>
    <h1>Ipsum!</h1>
    <p>Lorem ipsum dolor sit amet, consectetur adipiscing elit.
        Pellentesque tincidunt libero diam, nec imperdiet libero
        sodales quis. Nulla in pulvinar sem. Vivamus placerat
        ullamcorper sagittis. Proin varius, erat sed egestas semper,
        enim lectus viverra diam, id placerat est augue et turpis.
    </p>
</body>
</html>
```

Try it on your own code, and see if you can open the `index.html` and `ipsum.html` files.

404 - Not found

We're not done yet! This Is Our Bullshit Clickbait Version Of What Happens if we try to request a file that does not exist, such as <http://localhost:8000/hello.html>:



```
traceback (most recent call last):
  File "httpserver.py", line 36, in <module>
    fin = open('htdocs' + filename)
FileNotFoundError: [Errno 2] No such file or directory: 'htdocs/hello.html'
```

The server crashes and exits!

We need to catch the exception and return a 404 response:

```
while True:
    # Wait for client connections
    (...)

    # Get the client request
    (...)

    # Parse HTTP headers
    headers = request.split('\n')
    filename = headers[0].split()[1]

    # Get the content of the file
    if filename == '/':
        filename = '/index.html'

    try:
        fin = open('htdocs' + filename)
        content = fin.read()
        fin.close()

        response = 'HTTP/1.0 200 OK\n\n' + content
    except FileNotFoundError:

        response = 'HTTP/1.0 404 NOT FOUND\n\nFile Not Found'

    # Send HTTP response
    client.sendall(response)
```

client By using Codementor, you agree to our [Cookie Policy](#).

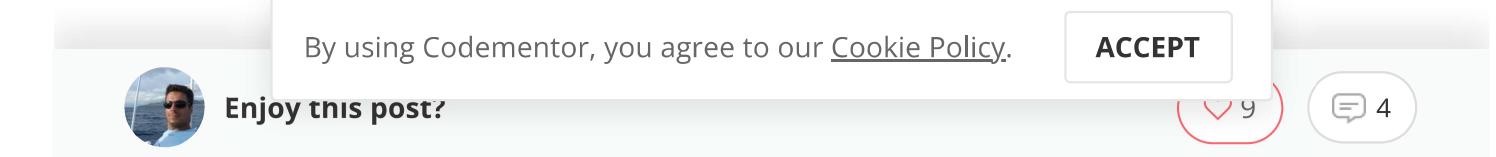
ACCEPT



Enjoy this post?



If you want, you can change the body of the Http 404 response to have personalized error messages..



The entire source code for this example can be found in [this gist](#).

[Python](#)[Server side](#)[Distributed systems](#)

Enjoy this post? Give **João Ventura** a like if it's helpful.

 9 4 SHARE

João Ventura

Python enthusiast with 10 years of experience!

I'm João Ventura, a software engineer from Portugal. Currently I work as a Computer Science Professor at the Polytechnique Institute of Setúbal, where I teach the Operating Systems and Distributed Computing courses. Some of my ar...

[FOLLOW](#) 4 Replies Leave a reply João Carlos 3 months ago

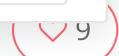
Hello João, very nice example where I could follow and make some adaptations in the socketser module in python. Now I have a question where I need to know how to send an image (.jpg, jpeg, png ...) to the client side, would you be able to exemplify ?



By using Codementor, you agree to our [Cookie Policy](#).

[ACCEPT](#)

Enjoy this post?

 9 4

 **avensis david** 10 months ago

very interesting thanks for sharing

<https://getappvalley.com/> <https://tutuappx.com/> <https://tweakbox.mobi/>

 Reply **Muhammad Abdullah** a year ago

are you

 Reply Show more replies

Find a Pair Programming Partner on Codementor

Want to improve your programming skills? Choose from 10,000+ mentors to pair program with.

 GET STARTED

Enjoy this post?

By using Codementor, you agree to our [Cookie Policy](#).

 ACCEPT